

CSE 142 Assignment 4, Fall 2023

4 Questions, 100 pts, due: 23:59 pm, Nov 22th, 2023

Your name: Airi Kokuryo

Student ID:

2086695

Instruction

- Submit your assignments onto **Gradescope** by the due date. Upload a `zip` file containing:

(1) The saved/latest `.ipynb` file, please **rename this file with your name included**.

(2) Also save your file into a pdf version, if error appears, save an html version instead (easy to grade for written questions).

For assignment related questions, please reach TA or grader through Slack/Piazza/Email.

- This is an **individual** assignment. All help from others (from the web, books other than text, or people other than the TA or instructor) must be clearly acknowledged.

Objective

- **Task 1:** EM algorithm (Mathematical Derivation) - **Optional Exercise**
- **Task 2:** K-Means implementation (Coding)
- **Task 3:** Kernel Methods with Noisy Setting (Coding)

Question 1. (Bonus OPTIONAL) EM algorithm, 20 pts

Derive the E-step and M-step update equations of EM algorithm for estimating the Gaussian mixture model $p(X; \theta) = \sum_{k=1}^K \pi_k N(x; \mu_k, \sigma_k^2)$ where π_k is the mixture weight with $\pi_k \geq 0$ and $\sum_{k=1}^K \pi_k = 1$, and μ_k, σ_k^2 are the mean and variance of the gaussian distribution corresponding to cluster k .

For the E-step, first prove that $z_{ik} = P(z_i = k | X, \mu, \sigma, \pi) = \frac{\pi_k N(x_i; \mu_k, \sigma_k^2)}{\sum_{k=1}^K \pi_k N(x_i; \mu_k, \sigma_k^2)}$. Then, for the M-step, show the derivation to compute the updates for (μ_k, π_k) . Note that, you don't need

to show the derivation for σ_k . For each derivation step, mention the concept applied (e.g. just 2-3 keywords, e.g. formula for expectation, independence of datapoints, $(f+g)' = f' + g'$, etc ...).

Hint: For the M-step, you need to solve for

$$\mu_k^t = \underset{\mu_k}{\operatorname{argmax}} E_{p(Z|X, \mu^{(t-1)}, \sigma^{(t-1)}, \pi^{(t-1)})} [\log p(X, Z | \mu, \sigma, \pi)] \text{ (and similarly for } \pi_k \text{) by applying the}$$

first order conditions for function optimization (take derivative and set it to zero). Note that the term $p(Z|X, \mu^{(t-1)}, \sigma^{(t-1)}, \pi^{(t-1)})$ is the one computed in the E-step, and uses fixed values for μ, σ, π from the previous iteration (t-1).

This exercise is optional, any effort will be rewarded with extra points on the assignment as a whole!

```
In [ ]: from IPython.display import Image
# Replace the figure name
Image(filename='question1_1.jpg')
```

Out []:

E-step & M-step

Gaussian Mixture Model

$$p(x|\theta) = \sum_{k=1}^K \pi_k N(x; \mu_k, \sigma_k^2)$$

 π_k = mixture weight w/ $\pi_k \geq 0$ & $\sum_{k=1}^K \pi_k = 1$
 μ_k : mean σ_k^2 : Variance

1. E-step

$$z_{ik} = p(z_i = k | x_i, \mu, \sigma, \pi)$$

$$= \frac{p(x_i | z_i = k, \mu, \sigma) p(z_i = k | \pi)}{p(x_i | \mu, \sigma)}$$

From the formula of marginal probability,

$$= \frac{p(x_i | z_i = k, \mu, \sigma) p(z_i = k | \pi)}{\sum_{k=1}^K p(x_i | z_i = k, \mu, \sigma) p(z_i = k | \pi)}$$

 where $p(x_i | z_i = k, \mu, \sigma) = N(x_i; \mu_k, \sigma_k^2)$
 and $p(z_i = k | \pi) = \pi_k$

therefore

$$= \frac{\pi_k N(x_i; \mu_k, \sigma_k^2)}{\sum_{k=1}^K \pi_k N(x_i; \mu_k, \sigma_k^2)}$$

2. M-step

Assume $\theta = \{\mu, \sigma, \pi\}$ the likelihood of X is given by

$$p(X|\theta) = \prod_{i=1}^N p(x_i|\theta) = \prod_{i=1}^N \prod_{k=1}^K \pi_k N(x_i; \mu_k, \sigma_k^2)$$

if the loss function is L ,

$$L = \log p(X|\theta) = \log \left\{ \prod_{i=1}^N \sum_{k=1}^K \pi_k N(x_i; \mu_k, \sigma_k^2) \right\}$$

$$= \sum_{i=1}^N \log \left\{ \sum_{k=1}^K \pi_k N(x_i; \mu_k, \sigma_k^2) \right\}$$

(i) For μ_k

$$\frac{\partial L}{\partial \mu_k} = \sum_{i=1}^N \frac{\partial}{\partial \mu_k} \log \left\{ \sum_{k=1}^K \pi_k N(x_i; \mu_k, \sigma_k^2) \right\}$$

$$= \sum_{i=1}^N \frac{\partial}{\partial \mu_k} \sum_{k=1}^K z_{ik} \left\{ \log \pi_k + \log N(x_i; \mu_k, \sigma_k^2) \right\}$$

$$= \sum_{i=1}^N \frac{\frac{\partial}{\partial \mu_k} \left\{ \pi_k N(x_i; \mu_k, \sigma_k^2) \right\}}{\sum_{k=1}^K \pi_k N(x_i; \mu_k, \sigma_k^2)}$$

$$= \frac{\sum_{i=1}^N \pi_k \frac{\partial}{\partial \mu_k} \left[\frac{1}{\sqrt{2\pi} \sigma_k} \exp \left\{ -\frac{(x_i - \mu_k)^2}{2\sigma_k^2} \right\} \right]}{\sum_{k=1}^K \pi_k N(x_i; \mu_k, \sigma_k^2)}$$

$$= \frac{\sum_{i=1}^N \pi_k \frac{1}{\sqrt{2\pi} \sigma_k} \exp \left\{ -\frac{(x_i - \mu_k)^2}{2\sigma_k^2} \right\} \left(-\frac{x_i - \mu_k}{\sigma_k^2} \right)}{\sum_{k=1}^K \pi_k N(x_i; \mu_k, \sigma_k^2)}$$

$$= \frac{\sum_{i=1}^N \pi_k N(x_i; \mu_k, \sigma_k^2) \left(-\frac{x_i - \mu_k}{\sigma_k^2} \right)}{\sum_{k=1}^K \pi_k N(x_i; \mu_k, \sigma_k^2)}$$

$$\therefore \frac{\partial L}{\partial \mu_k} = \sum_{i=1}^N z_{ik} \left(-\frac{x_i - \mu_k}{\sigma_k^2} \right)$$

$$\left(\because z_{ik} = \frac{\pi_k N(x_i; \mu_k, \sigma_k^2)}{\sum_{k=1}^K \pi_k N(x_i; \mu_k, \sigma_k^2)} \right)$$

From $\frac{\partial L}{\partial \mu_k} = 0$,

$$\mu_k = \frac{\sum_{i=1}^N z_{ik} x_i}{\sum_{i=1}^N z_{ik}} \quad (\text{update function})$$

Made with Goodnotes

In []: Image(filename='question1_2.jpg')

Out[]:

(ii) For π_k

$$\max_{\pi_k} L \quad \text{subject to} \quad \sum_{k=1}^K \pi_k = 1 \quad (\text{constraint})$$

$$\leftrightarrow \max_{\pi_k} Q = L + \lambda \left(\sum_{k=1}^K \pi_k - 1 \right) \quad (\lambda \in \mathbb{R})$$

$$\frac{\partial Q}{\partial \pi_k} = \sum_{i=1}^N \frac{\partial}{\partial \pi_k} \left\{ \log \left(\sum_{k=1}^K \pi_k N(x_i; \mu_k, \sigma_k^2) \right) \right\} + \lambda$$

$$= \sum_{i=1}^N \frac{N(x_i; \mu_k, \sigma_k^2) \pi_k}{\sum_{k=1}^K \pi_k N(x_i; \mu_k, \sigma_k^2) \pi_k} + \lambda$$

$$= \sum_{i=1}^N \frac{z_{ik}}{\pi_k} + \lambda = 0$$

$$\therefore \pi_k = - \frac{\sum_{i=1}^N z_{ik}}{\lambda}$$

Take the summation

$$\sum_{i=1}^N z_{ik} = -\lambda \pi_k$$

$$\sum_{k=1}^K (-\lambda \pi_k)$$

$$= -\lambda$$

$$= \sum_{k=1}^K \sum_{i=1}^N z_{ik}$$

$$= \sum_{i=1}^N 1 = N \quad \left(\sum_{k=1}^K \pi_k = 1, \sum_{k=1}^K z_{ik} = 1 \right)$$

Since $-\lambda = N$,

$$\pi_k = \frac{\sum_{i=1}^N z_{ik}}{N} \quad (\text{update function})$$

Made with Goodnotes

Question 2. (K-Means implementation, 25 pts)

Question 2.1. Implement K-means in Python from scratch. Complete following sub-functions `update_centroids` and `update_assignments`.

```
In [ ]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import make_blobs
```

```
In [ ]: def update_assignments(data, centroids):

    #####
    #### YOUR CODE HERE ####
    ## you will get cluster#
    ##assignments here #####
    k = np.shape(centroids)[0]
    assignment_list = []
    for i in range(k):
        #calculate norm of the vector
        norm = np.linalg.norm(data-centroids[i], axis=1)
        assignment_list.append(norm)
    #minimum element of the array in the particular axis
    assignments = np.argmin(assignment_list, axis=0)
    #####
    return assignments

def update_centroids(data, centroids, assignments):

    #####
    #### YOUR CODE HERE ####
    k = np.shape(centroids)[0]
    new_centroids = []
    for i in range(k):
        new_centroids.append(data[assignments == i].mean(axis=0))
    new_centroids = np.array(new_centroids)
    #####
    return new_centroids

def kmeans(data, centroids, max_iterations):

    for j in range(max_iterations):
        # update cluter assignments
        assignments = update_assignments(data,centroids)    # WRITE CODE FOR update_ass

        # update centroid locations
        centroids = update_centroids(data,centroids,assignments)  # WRITE CODE FOR upda

    # final assignment update
    assignments = update_assignments(data,centroids)
    return centroids, assignments
```

Question 2.2. Run your code on following toy dataset for different k-values, where $k = \{2, 3, 4, 6, 10\}$ and plot the cluster assignments for different k's as shown in following diagram.

```
In [ ]: from sklearn.datasets import make_blobs
import matplotlib.pyplot as plt

# Generate sample data
n_samples = 4000
n_components = 4

X, y_true = make_blobs(
```

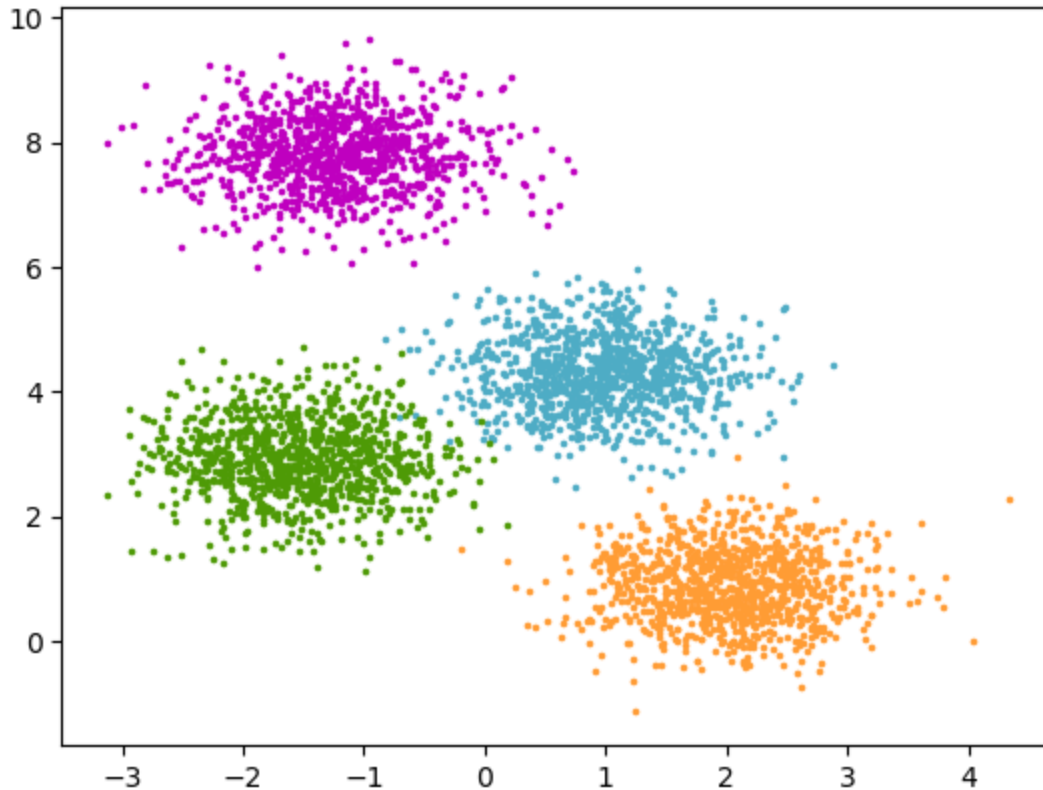
```

n_samples=n_samples, centers=n_components, cluster_std=0.60, random_state=0
)

colors = ["#4EACC5", "#FF9C34", "#4E9A06", "m"]

for k, col in enumerate(colors):
    cluster_data = y_true == k
    plt.scatter(X[cluster_data, 0], X[cluster_data, 1], c=col, marker=".", s=10)

```



```

In [ ]: import numpy as np

# function to get initial cluster centroids; we randomly choose k points from the dataset
def get_initial_clusters(k, X):
    random_indices = np.random.randint(0, X.shape[0], k)
    initial_centroids = X[random_indices]
    return initial_centroids
# your code here.
k_values = [2, 3, 4, 6, 10]

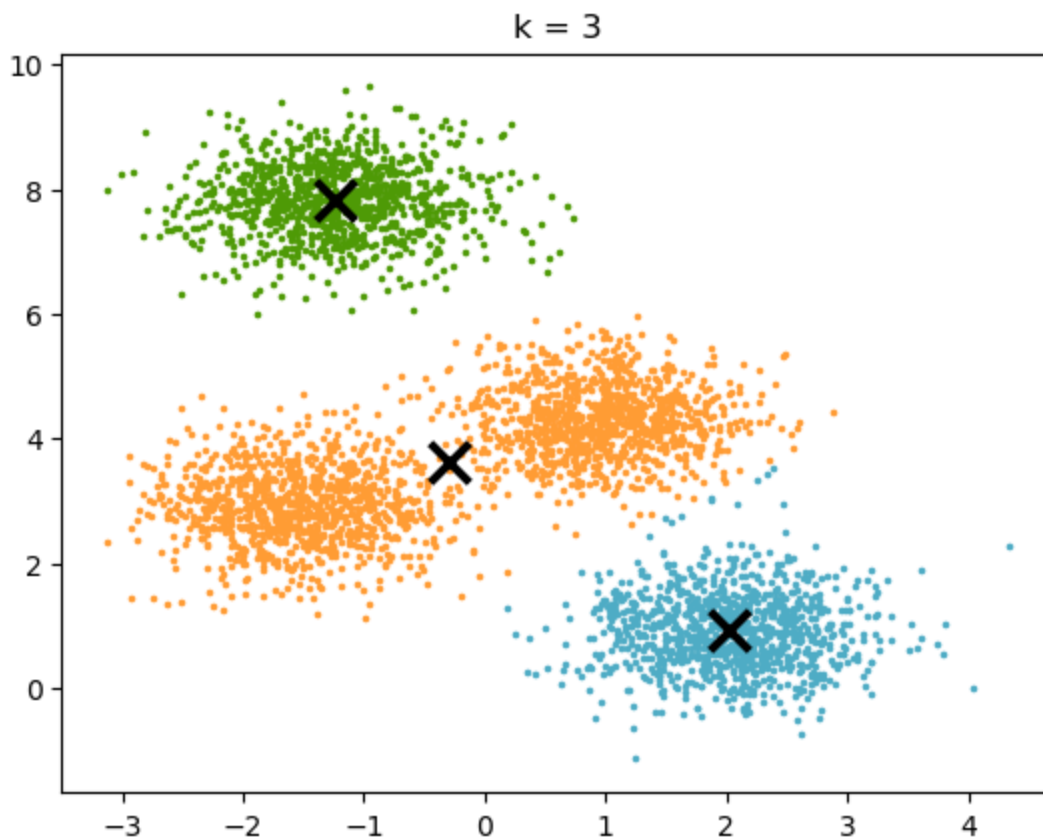
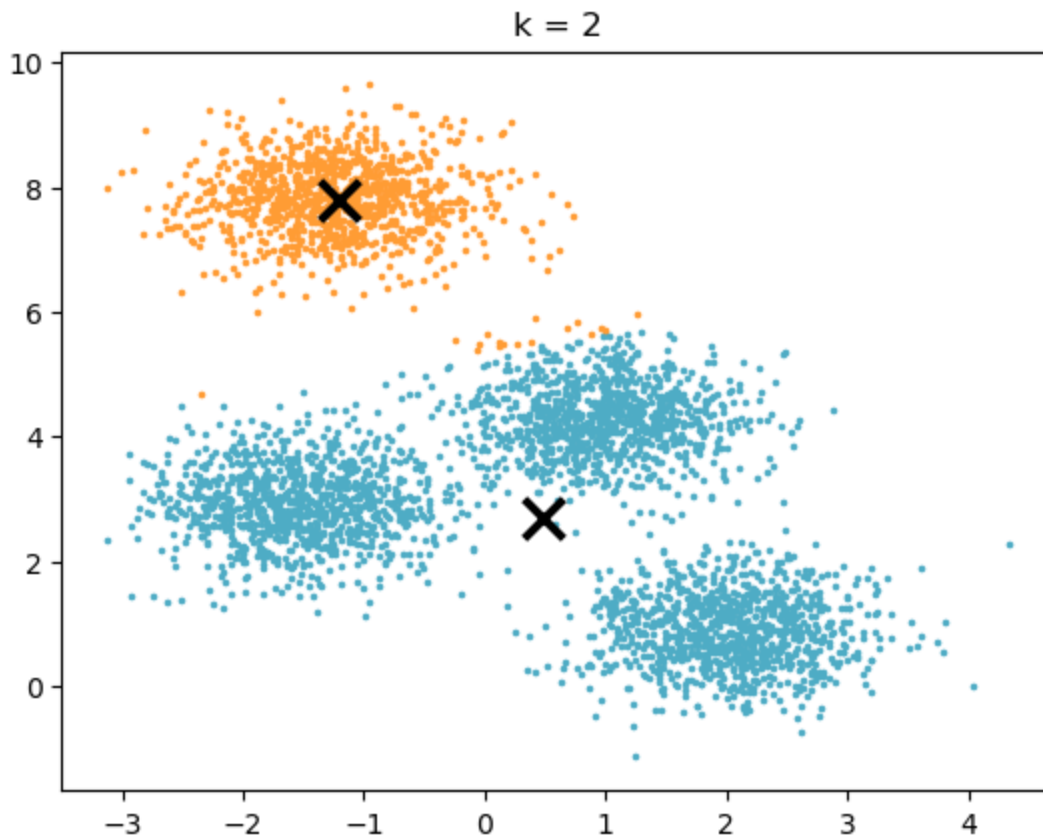
for k in k_values:
    # initialize centroids
    initial_centroids = get_initial_clusters(k, X)
    # run k-means
    final_centroids, assignments = kmeans(X, initial_centroids, max_iterations = 300)

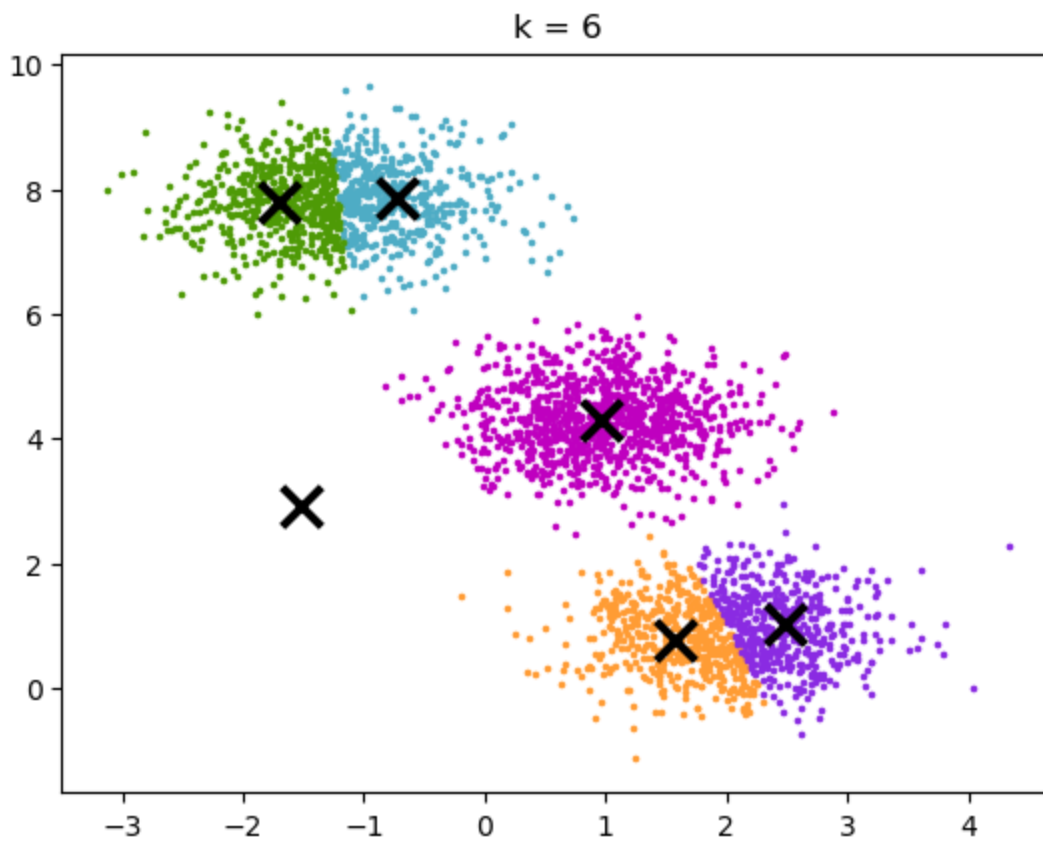
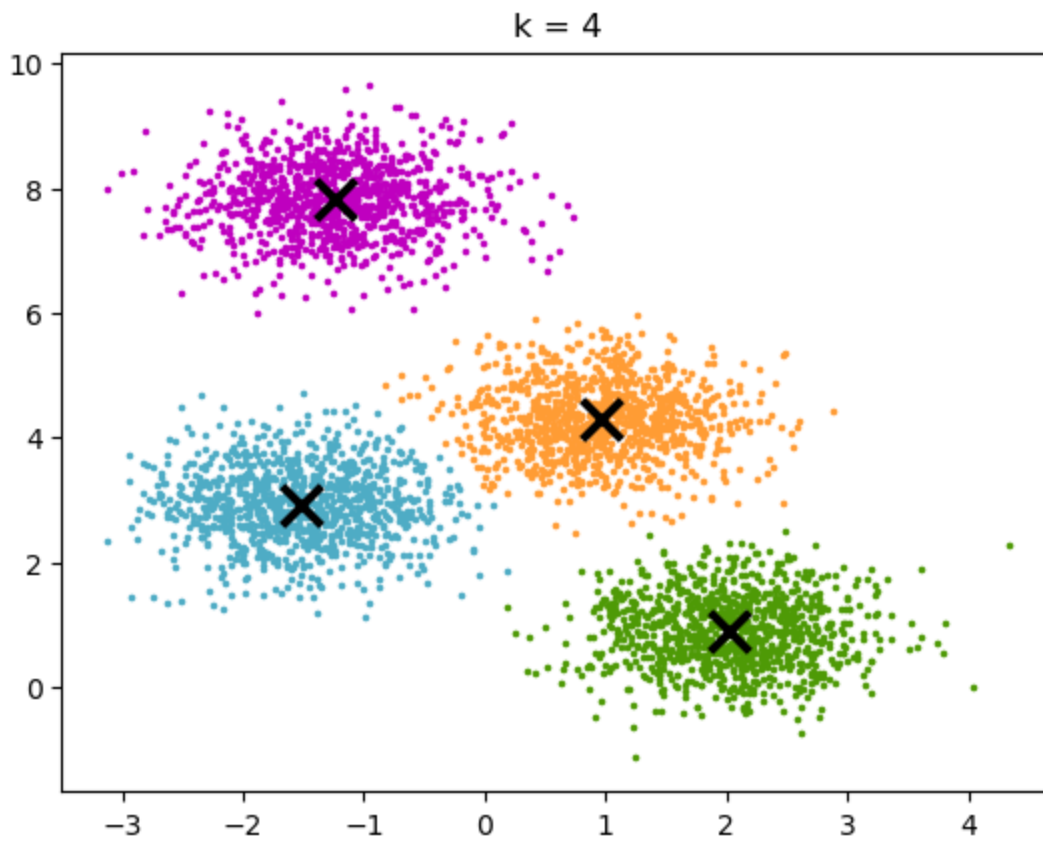
    # plot the results
    colors = ["#4EACC5", "#FF9C34", "#4E9A06", "m", "#8A2BE2"]
    plt.figure()
    for i, col in enumerate(colors[:k]):
        cluster_data = assignments == i
        plt.scatter(X[cluster_data, 0], X[cluster_data, 1], c=col, marker=".", s=10)

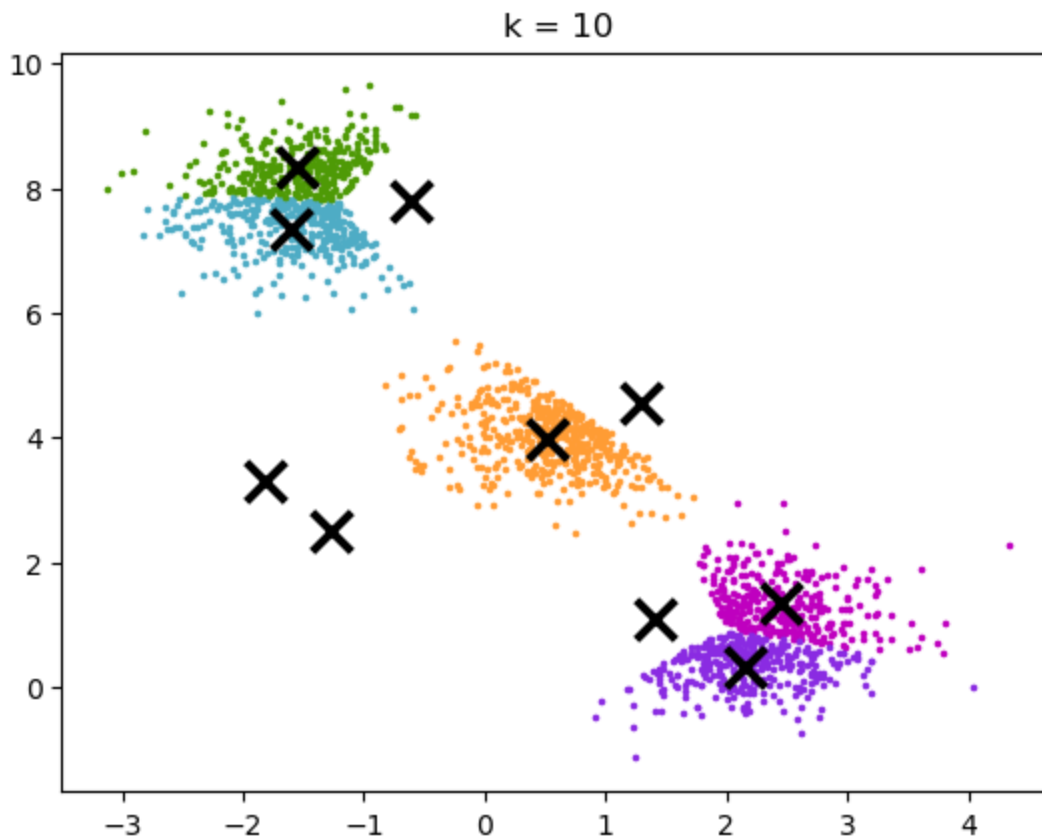
    # plot the centroids
    plt.scatter(final_centroids[:, 0], final_centroids[:, 1], marker="x", s=200, linewidth=2)

```

```
plt.title("k = {}".format(k))  
plt.show()
```







Question 3. (Kernel Methods with Noisy Setting, 75 pts)

SVM on synthetic dataset generated as follows:

- Draw 1000 (x_0, x_1) feature vectors from the 2-D Gaussian distribution with mean $\mu_+ = (1, 1)$ and $\Sigma_+ = \begin{bmatrix} 1, 0 \\ 0, 1 \end{bmatrix}$ and label them as $+1$.
- Draw 1000 (x_0, x_1) feature vectors from the 2-D Gaussian distribution with mean $\mu_- = (-1, -1)$ and $\Sigma_- = \begin{bmatrix} 3, 0 \\ 0, 3 \end{bmatrix}$ and label them as -1 .
- This gives you a 2000 example training set. Repeat the above to draw a test set the same way.

Use a SVM package (`scikit-learn svm.SVC` class) to learn SVMs with a variety of parameter settings.

(a -- 25 pts)

- Use an RBF kernel with parameters $C = 1, \gamma = 0.01$.
- For each training data with $+1$ label, randomly flip their label to -1 with probability **0.35**.
- For each training data with -1 label, randomly flip their label to $+1$ with probability **0.20**.

- Train with the above noisy training examples.
- Random flipping introduces the randomness. You can repeat multiple times (e.g. 20) and then report the average accuracy on the testing dataset (clean) in the noise parameter setting.

```
In [ ]: # Your code here
import numpy as np
from sklearn import svm
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score

# generate dataset for X
def generate_dataset(seed_data):
    np.random.seed(seed_data)

    dataset_number = 1000
    mean_positive = np.array([1, 1])
    mean_negative = np.array([-1, -1])

    cov_positive = np.array([[1, 0], [0, 1]])
    cov_negative = np.array([[3, 0], [0, 3]])

    positive_dataset = np.random.multivariate_normal(mean_positive, cov_positive, dataset_number)
    negative_dataset = np.random.multivariate_normal(mean_negative, cov_negative, dataset_number)

    features = np.vstack((positive_dataset, negative_dataset))

    return features

# generate dataset for y
def generate_labels(seed_data, seed_flip, bool_flip):
    np.random.seed(seed_data)
    positive_labels = np.ones(1000)
    negative_labels = -1 * np.ones(1000)

    if bool_flip:
        positive_labels = flip_label(positive_labels, 0.35, seed_flip)
        negative_labels = flip_label(negative_labels, 0.20, seed_flip)

    labels = np.hstack((positive_labels, negative_labels))
    return labels

# flip labels under its probability
def flip_label(label, probability, seed_flip):
    np.random.seed(seed_flip)
    flipped_data = np.random.rand(1000) < probability
    label[flipped_data] = -1 * label[flipped_data]
    return label

# get accuracy
def get_accuracy(X_train, X_test, y_train, y_test):
    svm = SVC(kernel='rbf', C=1.0, gamma=0.01, random_state = 42)
    svm.fit(X_train, y_train)
    y_pred = svm.predict(X_test)

    accuracy = accuracy_score(y_test, y_pred)
```

```

    return accuracy

accuracies = []
for i in range(20):
    X_train = generate_dataset(42)
    X_test = generate_dataset(24)
    y_train = generate_labels(42, i, True)
    y_test = generate_labels(24, 42, False)
    accuracies.append(get_accuracy(X_train, X_test, y_train, y_test))

print(f"The average of the accuracy is {np.mean(accuracies):.3f}")

```

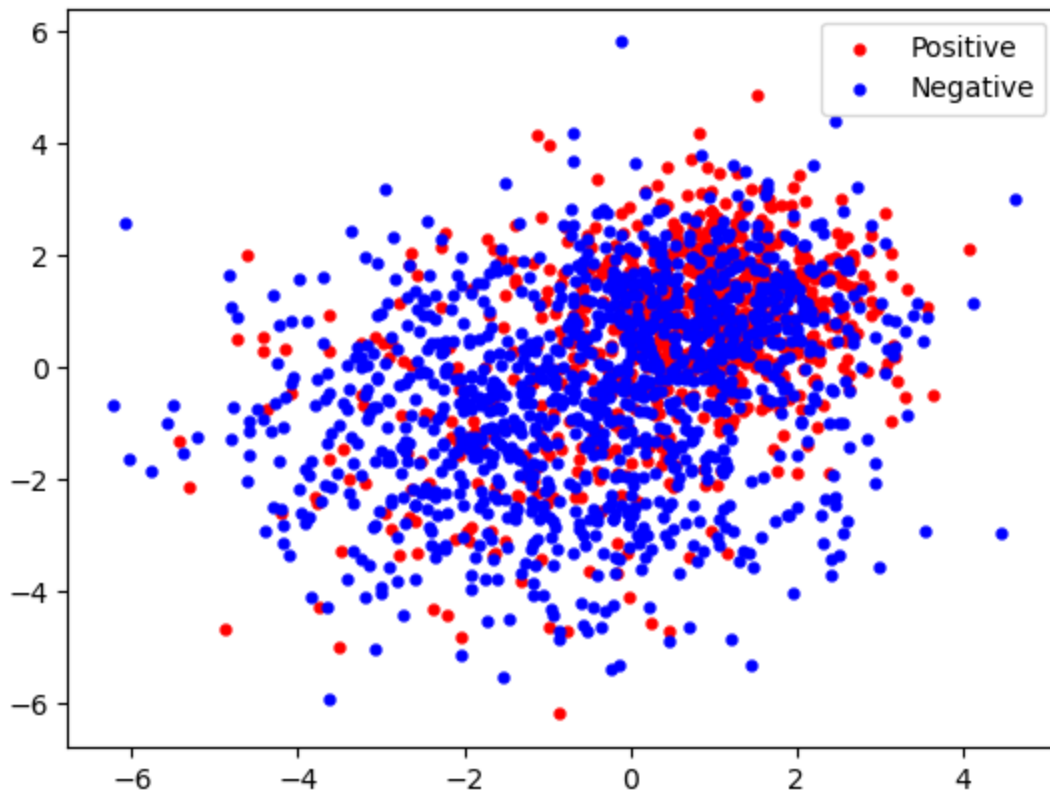
The average of the accuracy is 0.806

```

In [ ]: plt.scatter(X_train[y_train == 1, 0], X_train[y_train == 1, 1], c="r", marker=".", s=50)
plt.scatter(X_train[y_train == -1, 0], X_train[y_train == -1, 1], c="b", marker=".", s=50)
plt.legend()

```

Out[]: <matplotlib.legend.Legend at 0x16a00406110>



(b -- 25 pts) Open question

- Try using **K-Nearest Neighbors** to correct wrong labels before training.
- Then train the model with the newly processed training dataset.
- Report the accuracy on the testing dataset in the noise parameter setting. Do you observe performance improvement?

```

In [ ]: # Your code here

```

```

from sklearn.neighbors import KNeighborsClassifier

# Correct labels using k-Nearest Neighbors
def correct_labels_with_knn(X_train, y_train):
    knn = KNeighborsClassifier(n_neighbors = 5)
    knn.fit(X_train, y_train)
    corrected_labels = knn.predict(X_train)
    return corrected_labels

accuracies_with_knn = []

for i in range(20):
    X_train = generate_dataset(42)
    X_test = generate_dataset(24)
    y_train = generate_labels(42, i, True)
    y_test = generate_labels(24, 42, False)

    # Correct labels with knn
    y_train = correct_labels_with_knn(X_train, y_train)
    accuracies_with_knn.append(get_accuracy(X_train, X_test, y_train, y_test))

print(f"Accuracy after using KNN for label correction: {np.mean(accuracies_with_knn):.3}")

```

Accuracy after using KNN for label correction: 0.845

Performance improvement was observed by using the K-Nearest Neighbors to correct wrong labels before training.

(c -- 25 pts) Open question

- Try using **clustering (i.e., K-means, EM-clustering)** to correct wrong labels before training.
- Then train the model with the newly processed training dataset.
- Report the accuracy on the testing dataset in the noise parameter setting. Do you observe performance improvement?

```

In [ ]: # Your code here
from sklearn.cluster import KMeans

# Function to correct labels using K-means clustering
def correct_labels_with_kmeans(X_train):
    kmeans = KMeans(n_clusters=2, random_state=42, n_init="auto")
    kmeans.fit(X_train)
    y_train_corrected = kmeans.predict(X_train)
    return y_train_corrected

accuracies_with_kmeans = []

for i in range(20):
    X_train = generate_dataset(42)
    X_test = generate_dataset(24)
    y_train = generate_labels(42, i, True)
    y_test = generate_labels(24, 42, False)

```

```
# Correct labels with K-means
y_train = correct_labels_with_kmeans(X_train)
accuracies_with_kmeans.append(get_accuracy(X_train, X_test, y_train, y_test))

print(f"Accuracy with K-means corrected labels: {np.mean(accuracies_with_kmeans):.3f}")
```

Accuracy with K-means corrected labels: 0.017

The performance got worse when K-means was used before training.