

| | |
|------------------------|--|
| i.i.d | examples are drawn independently from the same distribution – and this same distribution is used to generate both the training set and the test data. |
| overfitting | when a model learns to replicate noise or spurious patterns in the training data. An overfit model may closely match training data but may fail to generalize to unseen test data. |
| underfitting | when a model is too simple, informed by too few features or regularized too much, and can not adequately capture the underlying structure of the data. |
| Supervised learning | Has a clear objective, and the model's performance can be quantitatively measured using metrics. Overfitting risks and labeled data may be time-consuming to obtain. |
| Unsupervised learning | Does not require labeled data, and may lead to insights that were not seen in supervised learning. Not always have a clear semantic meaning, and may not be suited for tasks that require precise categorization or prediction without labeled examples. |
| Reinforcement Learning | Agents and environments, rewards and punishments. ML paradigm where an agent learns to make decisions by interacting with an environment. |
| Neural Networks | A computational model inspired by the structure and functioning of the human brain. Perceptrons, activation functions, feedforward and backpropagation |
| Deep Learning | Convolutional Neural Networks, Recurrent Neural Networks: a subset of machine learning that involves the use of artificial neural networks with multiple layers |
| Evaluation Metrics | quantitative measures used to assess the performance of a machine learning model |
| Ensemble learning | Boosting(ada boost) |
| Bias and Variance | There is a trade-off between bias and variance, and finding the right balance is crucial for building models that generalize well. |
| Bias | High bias can lead the model to underfit the data, meaning it is too simplistic and fails to capture the underlying patterns in the training data. |
| Variance | High variance can lead to overfitting, where the model performs well on the training data but fails to generalize to new, unseen data. |
| Regularization | L1 and L2 regularization. a technique used in machine learning to prevent overfitting and improve the generalization performance of a model. |
| Clustering Algorithms | K-Means, hierarchical clustering. algorithms are used to group similar data points into clusters. Unsupervised learning |
| PCA(主成分分析) | Dimensionality reduction technique. The goal is to retain the most important information in the data while reducing its dimensionality. |
| Deployment | process of taking a trained machine-learning model and making it available for use in a real-world environment |

Ch. 2 Bayes' rule

$$P(x|y) = P(y|x)P(x) / P(y)$$

| | | | |
|-----------|--------|------------------------|--------|
| Posterior | P(x y) | Likelihood | P(y x) |
| Prior | P(x) | Normalization constant | P(y) |

Assume we have a HHHHTTHHT

- Maximum likelihood estimate for h: 5/8
- Prior density for the various values of h ∈ [0,1] give the formula for the posterior probability density p(h|HHHHTTHHT) using the prior p(h)

$$p(h|HHHHTTHHT) = \frac{p(HHHHTTHHT|h)p(h)}{p(HHHHTTHHT)}$$

Marginal likelihood: $p(HHHHTTHHT) = \int_0^1 q^5(1-q)^3 p(q) dq$

$$= \frac{h^5(1-h)^3 p(h)}{\int_0^1 q^5(1-q)^3 p(q) dq}$$

Conditional probability:

containing the word "payment", given that the email is a spam email, is 72%. Suppose that the conditional probability of an email being spam, given that it contains the word "payment", is 8%. Find the ratio of the probability that an email is spam to the probability that an email contains the word "payment"

payment = p(a), spam = p(b)
 $p(a|b) = 0.72 = p(a^*b)/p(b)$, $p(b|a) = 0.08 = p(a^*b)/p(a)$
 $p(b|a)/p(a|b) = p(b)/p(a) = 0.72/0.08 = 9/1$

There are two boxes. Box 1 contains three red and five white balls and box 2 contains two red and five white balls. A box is chosen at random $p(\text{box} = 1) = p(\text{box} = 2) = 0.5$ and a ball chosen at random from this box turns out to be red. What is the posterior probability that the red ball came from box 1?

$P(\text{box}=1|\text{red})=?$

$P(\text{red}|\text{box}1) = (3/8)/(1/2) = 3/4$

$P(\text{box}1|\text{red}) = p(\text{red}|\text{box}1)*p(\text{box}1)/p(\text{red}) = (3/4*1/2)/(37/56)$

Ch. 3 Regularized Least Squares

| | |
|---------------------------|--|
| Regularized least squares | Distance metrics using norm. L1 is $ x+y = 1$ (square), L2 is $\sqrt{x^2+y^2}$ (circle) |
| Curse of Dimensionality | The volume of space increases as the dimensionality increase |

Ch. 4 Loss Function, Linear classification, SGD

Linear classification: the goal is to find the linear decision boundary

| | |
|---------------|--|
| Loss function | Serves as a quantitative measure of how well the model is performing on a given task by comparing its predictions to the actual (ground truth) values. |
|---------------|--|

$$MSE = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

y_i : actual value for sample i, \hat{y}_i : predicted value

Classification loss functions: Binary Cross-Entropy Loss

$$-\frac{1}{N} \sum_{i=1}^N [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$$

| | |
|-------------------|--|
| Linear Regression | used to predict the value of a variable based on the value of another variable |
|-------------------|--|

Linear regression formula: $y_i = f(x_i, \beta) + e_i$

Ch. 5 Perceptron

| | |
|------------|--|
| perceptron | One of the simplest neural network architectures |
|------------|--|

Binary input: x, weight: w, bias: b

Linear decision boundary: $z = w \cdot x + b = 0$

Passes the weighted sum through an activation function to produce a binary output. Converges on linearly separable data.

| x_0 | x_1 | x_2 | y |
|-------|-------|-------|-----|
| -1 | 2 | 1 | -1 |
| 1 | 1 | 1 | +1 |
| -2 | 3 | 2 | -1 |
| 2 | -1 | -1 | +1 |

Initially, $w = (-1, 1, 1)$

Simulate one pass through the following data with the perceptron algorithm described in the lecture and homework. Start with $w = (-1, 1, 1)$ and show the resulting weight vector after each example. (Assume that the perceptron algorithm predicts incorrectly when $w \cdot x = 0$, and ignore the bias term.)

-first example: $w \cdot x = 4$, wrong,

$w = w + yx = (-1, 1, 1) - (-1, 2, 1) = (0, -1, 0)$

-second example $w \cdot x = -1$, wrong,

$w = w + yx = (0, -1, 0) + (1, 1, 1) = (1, 0, 1)$

-third example $w \cdot x = 0$,

wrong, $w = w + yx = (1, 0, 1) - (-2, 3, 2) = (3, -3, -1)$

-fourth example, $w \cdot x = 10$, correct no update.

Ch. 6 SVMs (support vector machines)

The support vectors are the closest points to the boundaries. Are the points that lie on the margin boundaries $wx+b=1$ and $wx+b=-1$

| | |
|-------------------------------|--|
| SVM and perceptron difference | while both perceptrons and SVMs are linear classifiers, SVMs emphasize maximizing the margin and provide more flexibility through the use of kernels for handling nonlinear separation |
| SVM | Classification algorithm that finds the maximum margin separating hyperplane |
| Optimization for SVM | Maximizing γ . Constraints will be that all training instances are correctly classified. |
| Properties | w is a linear combination of the support vectors/ pos and neg examples contribute equally to w |
| Hard margin | Hard constraint \uparrow w (and b) depend only on support vectors |
| Soft margin | Allow it to make some errors but it will have to pay a price for each error...this price is slack |

Ch. 7 Kernels

$$x = \langle x_1, x_2 \rangle \rightarrow \Phi(x) = \langle x_1^2, \sqrt{2}x_1x_2, x_2^2 \rangle$$

$$\Phi(x) * \Phi(z) = (x * z)^2$$

This is calculated by inserting the method inside the given function

Kernelized SVM Prediction: (Duality)

Functional margin: $y(w^*x + b)$

Geometric margin: $y(w^*x + b)/||w||_2 \rightarrow$ maximizing this is equiv to

$\min_{w,b} \|w\|_2$ subject to $y_i(w * x_i + b) \geq 1$ for all i , and
 $\min_{w,b} \frac{1}{2}(w * w)$ subject to $1 - y_i(w * x_i + b) \leq 0$ for all i
 $w^* = \sum_{i=1}^n \alpha_i^* y_i x_i$ means w is a weighted sum of examples (if without $*$)

$$b = -\frac{1}{2} \left(\min_{i: y_i = +1} (w^* * x_i) + \max_{j: y_j = -1} (w^* * x_j) \right)$$

Support vectors: $(1,0), (0,1)$
Equation for maximizing margin: $y=wx + b, w=-1$ so $x+y-b=0$
Therefore $|x+y-b|/\sqrt{(1+1)}, b=3/2 \rightarrow y = -x + 3/2$
Geometric margin: $|2-2/3|/\sqrt{(2)} = \sqrt{(2)}/4$

| | |
|--------------------|---|
| Mercer's condition | ensures the positive definiteness of a kernel function. a necessary and sufficient condition for a function to be a valid kernel. |
|--------------------|---|

Ch. 8 Naïve Bayes

| | |
|-------------|---|
| Naïve Bayes | Assume that the features are conditionally independent given the class. |
|-------------|---|

| x_1 | x_2 | y (label) | $p(x_1, x_2 y) = p(x_1 y) p(x_2 y)$ |
|-------|-------|-------------|---|
| 1 | 1 | + | Naïve Bayes estimates: $p(y = +) = 3/7; p(y = -) = 4/7;$ $p(x_1=1 y=+) = 2/3; p(x_1=1 y=-) = 1/2;$ $p(x_1=0 y=+) = 1/3; p(x_1=0 y=-) = 1/2;$ $p(x_1=1, x_2=0 y=+) = 2/3 \cdot 1/3 = 2/9$ $p(x_1=1, x_2=0 y=-) = 1/2 \cdot 1/2 = 1/4$ Thus Naïve Bayes estimates $p(+ x) \propto 3/7 \cdot 2/9 = 2/21$, and $p(- x) \propto 4/7 \cdot 1/4 = 3/21$ |
| 1 | 0 | + | |
| 0 | 1 | + | |
| 0 | 0 | - | |
| 0 | 0 | - | |
| 1 | 1 | - | |
| 1 | 1 | - | |

Ch. 9 K-Nearest-Neighbors

Is a non-parametric method. Normalize dimensions (distance metric)
Decision boundaries change with k (ex, change distance metrics)
Noise can cause problems; 10 % chance of fail

Ch. 10 Unsupervised learning (k-means)

- Steps to finding k
- Randomly initialize center for a k th cluster
 - assign example n to the closest center
 - points assigned to cluster k
 - re-estimate center of cluster k
 - return cluster assignments

Ch. 11 Gaussian Distribution, Expectation Maximization

| | |
|-------------|---|
| E step (P2) | Find the expectation of where the new point would be assigned to $\mathbb{P}(\mathbf{Z} \mathbf{X}, \pi^{\text{current}})$ |
| M step (P1) | Change clusters to a different location $\pi^{\text{new}} = \operatorname{argmax}_{\pi} \sum_{\mathbf{Z}} \mathbb{P}(\mathbf{Z} \mathbf{X}, \pi^{\text{current}}) \cdot \ln \mathbb{P}(\mathbf{Z}, \mathbf{X} \pi)$ |

Ch. 12 Supervised learning (Entropy)

| | |
|---------------|---------------------------------------|
| Entropy | An amount of information in something |
| Decision Tree | Small and simple trees are better |

$$H(s) = -p_+ \log(p_+) - p_- \log(p_-)$$

Ch. 13 Reinforcement Learning: Markov Decision Processes (MDPs)

| | |
|---------------------|---|
| Utilities | How much it is worth can be different in people; additive or discounted utility |
| Optimal utility | Bellman Equations; values change over time but the rewards don't change. |
| ModelBased Learning | algorithm builds a model of the underlying structure or patterns in the data during the training process |
| Model Free learning | learning optimal strategies or policies directly from observed data, rather than creating an internal representation of the system. |

Optimal value function

$$V^*(s) = \max_a Q^*(s, a)$$

Optimal value action function

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

Optimal policy

$$\pi^*(s) = \operatorname{argmax}_a Q^*(s, a)$$

Value iteration update equations


$$V_{i+1}(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_i(s')]$$

- Transition probability function: $T(s, a, s')$
Probability of transitioning from s to s' under a
- Reward function: $R(s, a, s')$
Immediate reward obtained when transitioning from s to s' by a

- Discount factor: γ
Parameter to determine the importance of future rewards in decision making process
- Action taken by an agent in a given state: a
- Current state: s
- Q-learning: $Q(s, a) = (1 - \alpha) \cdot Q(s, a) + \alpha \cdot (R(s, a, s') + \gamma \max_a Q(s', a))$

Ch. 14 Ensemble Methods and Boosting


| | |
|------------------|---|
| Ensemble Methods | avoid overfitting, increase weights of the misclassified dots |
| AdaBoost | Give misclassified instances a higher weight. Add up the weak classifiers |



Ch. 15 Deep Learning

Multilayer perceptron: neurons are arranged in layers (input, hidden, output)

| | |
|-----------------|---|
| Backpropagation | update parameters by propagating the error backward through the network |
|-----------------|---|



$$\frac{dL}{dW_2} = \frac{dL}{dz_1} * \frac{dz_1}{dw_2}$$

$$\delta_1 = \frac{dL}{da_1} = \frac{dL}{dz_1} * \frac{dz_1}{da_1} = \frac{dL}{dz_1} * f'(a_1)(1 - f(a_1))$$

$$\delta_2 = \frac{dL}{da_2} = \frac{dL}{dz_1} * \frac{dz_1}{dz_2} * \frac{dz_2}{da_2} = f'(a_2) * w_2^1 * \delta_1$$

$$\delta_5 = \delta_2 * w_5^2 + \delta_3 * w_5^3 + \delta_4 * w_5^4$$

$$a_2 = w_5^2 * z_5, z_2 = f(a_2), a_1 = w_2^1 \operatorname{sigmoid}(z_2) + \dots + w_4^1 \operatorname{sigmoid}(z_4)$$

$$z = \sum_i \operatorname{sigmoid}(a_i) \text{ where } i = [2, 3, 4]$$

$$z = \operatorname{sigmoid}(w_2^1 \operatorname{sigmoid}(f(w_5^2 * z_5)) + \dots + w_4^1 \operatorname{sigmoid}(f(w_5^4 * z_5)))$$